Implementing deep neural networks for classification and regression tasks

Alireza Afzal Aghaei[1]*

Abstract

In this report, we explain some techniques in deep neural networks by starting from multi-layer perceptron (MLP) model and it's mathematical definitions, then discuss convolutional neural networks (CNNs) which are used for image classification and finally review recurrent neural networks (RNN) as a powerful mathematical tool for time series forecasting. To increase generalization of the models we explained some regularization methods such as weight regularization and dropout. To show the accuracy of models we use some benchmark datasets for classification tasks. First, the iris dataset is applied on a simple MLP neural network, then we used the CNN model on mnist and fashion mnist dataset and finally applied RNNs on a sentiment analysis task.

Keywords

Deep learning, Convolutional neural networks, Recurrent neural networks

[1]*Department of Computer Science, Shahid Beheshti University, Tehran, Iran*
Corresponding author: alirezaafzalaghaei@gmail.com

Contents

1. Introduction

ANNs began as an attempt to exploit the architecture of the human brain to perform tasks that conventional algorithms had little success with. They retained the biological concept of artificial neurons, which receive input, combine the input with their internal state (activation) and an optional threshold using an activation function, and produce output using an output function. The initial inputs are external data, such as images and documents. The ultimate outputs accomplish the task, such as recognizing an object in an image. The important characteristic of the activation function is that it provides a smooth transition as input values change, i.e. a small change in input produces a small change in output. The network consists of connections, each connection providing the output of one neuron as an input to another neuron. Each connection is assigned a weight that represents its relative importance. A given neuron can have multiple inputs and output connections. A biological neuron made of a cell body (Soma), dendrites and an axon. Dendrite receives signals from other neurons, Soma sums all the incoming signals to generate input, Axon fires when the sum reaches a threshold value and sends a

Figure 1. A biological neuron
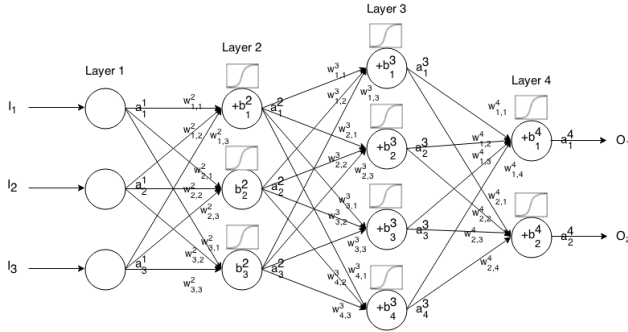

Figure 2. An artificial neuron

Figure 3. An Artificial neural network

signal to the synapse where is the point of interconnection of one neuron with other neurons.

Artificial neuron simulates the structure of a biological neuron. As of the simplicity of derivative calculation, it uses the sum of the product of input and corresponding weights. It fires if the summation reaches to a threshold. Note that, because of the good approximation of fuzzy activation functions, the function f may be more complex than a simple threshold function. These functions should be continuous and (weakly) differentiable. For example, *Tanh*, *Relu* and its variants, *Sigmoid* and radial basis functions (*RBFs*) are some of them. Same as linear regression, by finding a set of appropriate weights, we can solve a linear regression problem with this approach.

An artificial neural network (ANN) consists of many single neurons connected with a specific structure. Multi-layer perceptron (MLP) is the most basic network in which neurons are stacked in several layers. Convolutional and recurrent neural networks are other types of ANNs. By training a network

we can find optimal weights which leads to a good approximation of the desired dataset. This process can be divided into three major phases: Forward pass, Backward pass and, Updating weights. In the two next Sections, we explain the mathematics behind state-of-the-art networks such as MLP and CNN which can be trained using a backpropagation algorithm and gradient descent optimizer for supervised tasks.

2. Multi layer perceptron (MLP)

In this section, we explain the backpropagation algorithm for MLP networks and introduce some activation and loss functions that are used for classification and regression tasks.

2.1 Forward phase

Suppose a d-dimensional dataset contains n different samples named $X^{n \times d}$ and k-dimensional target values $Y^{n \times k}$. By defining a weight matrix $W_1^{d \times h1}$ and using matrix multiplication we have

$$z^{(1)} = X \cdot W_1 \tag{1}$$

which $Z^{(1)}$ denotes first hidden layer of network. This process is equal to apply simple weighted summation for whole dataset. By applying a nonlinear activation function we have

$$a^{(1)} = \sigma(z^{(1)}). \tag{2}$$

It's clear that the shape of $a^{(1)}$ is $(n \times h1)$. We can repeat this process to build deeper network.

$$z^{(i)} = a^{(i-1)} \cdot W_i \tag{3}$$

$$a^{(i)} = \sigma(z^{(i)}) \tag{4}$$

for $i = 2, \ldots, r-1$. In the last layer, we have

$$z^{(r)} = a^{(r-1)} \cdot W_r \tag{5}$$

$$\hat{y} = \sigma(Z^{(r)}) \tag{6}$$

which $\hat{y}$ is the prediction of neural network w.r.t input data X and weights $\{W_i\}_{i=1}^r$.

2.2 Loss function

To see the accuracy of the network, we need a measure function to show the cost of prediction. The most known loss functions are mean squared error (MSE) end cross entropy (Xentropy) which are used in regression and classifications tasks, respectively. There are many other loss functions such as mean absolute error, Vapnik's ε-insensitive loss function, etc. Here we explain MSE loss function:

$$J_{MSE}(W) = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}_i)^2 \tag{7}$$

Figure 4. Different optimizers

Figure 5. Dropout

2.3 Backward phase: Backpropagation

To find optimal weights with mathematics optimization tools, we need the derivative of loss function w.r.t weights $W_{i,j}^r$. Here we provide a well known technique known as Backpropagation. by defining

$$\delta^{(l)} = -(y - a^{(l)}) \odot \sigma'(z^{(l)}) \tag{8}$$

For last layer and

$$\delta^{(l)} = (\delta^{(l+1)} \cdot W_l^\mathsf{T}) \odot \sigma'(z^{(l)}) \tag{9}$$

the derivatives can be computed by

$$\frac{\partial J(W)}{\partial W_l} = a^{(l)\mathsf{T}} \cdot \delta^{(l+1)} \tag{10}$$

For better explanation see Brilliant backpropagation tutorial

2.4 Update phase

Function minimization is an essential topic in mathematics. When the function $f(x)$ has quadratic form, there is a special optimizer that can rich global minimum efficiently. But for arbitrary nonlinear functions, we need to use some parametric iterative algorithms to reach a minimum point that maybe is a local. Based on the derivative order needed by algorithms the classified into two categories known as first and second-order methods. While first-order methods need to save gradient vector in memory, second-order optimizers save a matrix of derivatives. Since these methods are very memory intensive for large datasets, some algorithms are developed which approximate that matrix among gradient calculation. The most known algorithms are Limited memory-BFGS (L-BFGS). Because of simple implementation, here we explain the gradient descent algorithm which is a very basic first-order optimizer. By starting from an initial point w_0 and goal function L, we can reach a local minimum with this update rule:

$$w_i = w_i - \eta \frac{\partial L}{\partial w_i} \tag{11}$$

where η called step length or learning rate. Note that, by choosing an appropriate value for step length we can reach the global minimum. To accelerate the learning process, researchers extended this update rule by adding some terms or changing step length base on some conditions. Momentum, Nestrov, Adagrad, AdaDelta, RMSProp, Adam are some extensions of vanilla gradient descent. see ruder.io for better formulation and pseudo-codes. Because of the effect of momentum technique on accelerating the training the neural networks and escape from local minimums, we note this method.

$$v = \mu v - \eta \frac{\partial L}{\partial w_i}$$
$$w_i = w_i + v$$

where μ is a hyper-parameter which is usually set to 0.9.

3. Regularization

Several techniques are used to overcome the over-fitting problem. Among them, weight decay and dropout are major methods.

3.1 Weight regularization

In weight decay, we add a positive function of unknown weights to the loss function. This leads to small weights and helps us to prevent over-fitting. because of simplicity here we explain L2-regularization which is named Tikhonov regularization in mathematics. This regularization can be defined by adding

$$\alpha \sum_{i=1}^{r} \|W_i\|_F \tag{12}$$

term to the loss function, where $\|\cdot\|_F$ denotes Frobinious norm and α is the regularization coefficient.

3.2 Dropout

Dropout is a regularization technique for reducing overfitting in neural networks by preventing complex co-adaptations on

Figure 6. Convolution operator

Figure 7. Pooling layer

Figure 8. RNN

training data. It is a very efficient way of performing model averaging with neural networks. The term "dropout" refers to dropping out units (both hidden and visible) in a neural network.

4. Convolutional neural network (CNN)

In this section, we introduce convolution and pooling operators used in Convolutional neural networks. Then derive the backpropagation formulas for these operators.

4.1 Convolution block
The objective of the Convolution Operation is to extract the high-level features such as edges, from the input image. ConvNets need not be limited to only one Convolutional Layer. Conventionally, the first ConvLayer is responsible for capturing the Low-Level features such as edges, color, gradient orientation, etc. With added layers, the architecture adapts to the High-Level features as well, giving us a network that has the wholesome understanding of images in the dataset, similar to how we would. There are two types of results to the operation — one in which the convolved feature is reduced in dimensionality as compared to the input, and the other in which the dimensionality is either increased or remains the same. This is done by applying Valid Padding in case of the former, or Same Padding in the case of the latter.

4.2 Pooling block
Similar to the Convolutional Layer, the Pooling layer is responsible for reducing the spatial size of the Convolved Feature. This is to decrease the computational power required to process the data through dimensionality reduction. Furthermore, it is useful for extracting dominant features that are rotational and positional invariant, thus maintaining the process of effectively training of the model. There are two types

of Pooling: Max Pooling and Average Pooling. Max Pooling returns the maximum value from the portion of the image covered by the Kernel. On the other hand, Average Pooling returns the average of all the values from the portion of the image covered by the Kernel. Max Pooling also performs as a Noise Suppressant. It discards the noisy activations altogether and also performs de-noising along with dimensionality reduction. On the other hand, Average Pooling simply performs dimensionality reduction as a noise suppressing mechanism. Hence, we can say that Max Pooling performs a lot better than Average Pooling.

The Convolutional Layer and the Pooling Layer, together form the ith layer of a Convolutional Neural Network. Depending on the complexities in the images, the number of such layers may be increased for capturing low-levels details even further, but at the cost of more computational power. After going through the above process, we have successfully enabled the model to understand the features. Moving on, we are going to flatten the final output and feed it to a regular Neural Network for classification purposes.

4.3 Backward phase
See this and this for backward phase of convolution and pooling layers.

5. Recurrent neural network (RNN)

A recurrent neural network (RNN) is a class of artificial neural networks where connections between nodes form a directed graph along a temporal sequence. This allows it to exhibit temporal dynamic behavior. Unlike feedforward neural networks, RNNs can use their internal state (memory) to process sequences of inputs. This makes them applicable to tasks such as unsegmented, connected handwriting recognition or speech recognition.

5.1 LSTM
A common LSTM unit is composed of a cell, an input gate, an output gate and a forget gate. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.

5.2 GRU
The GRU is like a long short-term memory (LSTM) with forget gate[2] but has fewer parameters than LSTM, as it lacks an output gate. GRU's performance on certain tasks of polyphonic music modeling and speech signal modeling was

Figure 9. LSTM cell

Figure 10. GRU

found to be similar to that of LSTM. GRUs have been shown to exhibit even better performance on certain smaller datasets. However, as shown by Gail Weiss & Yoav Goldberg & Eran Yahav, the LSTM is "strictly stronger" than the GRU as it can easily perform unbounded counting, while the GRU cannot. That's why the GRU fails to learn simple languages that are learnable by the LSTM.

5.3 Bidirectional
Bidirectional Recurrent Neural Networks (BRNN) connect two hidden layers of opposite directions to the same output. With this form of generative deep learning, the output layer can get information from the past (backward) and future (forward) states simultaneously.

5.4 Backward phase
The error backpropagation algorithm for recurrent networks is named as backpropagation through time (BPTT). see this link for more information about BTPP.


Figure 11. Bidirectional RNN

6. Text analysis

6.1 Embedding
Word embedding is a technique used to represent documents with a dense vector representation. The vocabulary in these documents is mapped to real number vectors. Semantically similar words are mapped close to each other in the vector space.

7. Implementation

In all of the examples for MLP training, we first normalize input data by subtracting the mean of a column for each data sample then dividing it by variance of the column. In the regression datasets, we do the same scaling on target values, but in the classification tasks, one-hot encoding is used. This process is done automatically by following the implemented class.

7.1 MLP model
The implementation of our MLP class is defined as a simple class named MLP.

```
MLP(hidden_layer, activation,
epoch, eta, beta, alpha, mu,
batch_size, verbose, task)
```

where options are

- **hidden_layer**: A list where each item defines the number of neurons in the hidden layer. It's obvious that, the length of list denotes number of hidden layers.

- **activation**: A single class that inherits Activation class and applied to each layer. Default choices are Sigmoid, Tanh, ReLU and LeakyReLU.

- **epoch**: The number of iterations.

- **eta**: Step length or learning rate

- **beta**: A number which multiplied to initial weights, initial weights have normal distribution

- **alpha**: Regularization coefficient

- **mu**: Momentum parameter

- **batch_size**: The number of batch passed to batch gradient descent

- **verbose**: (optional) Use this to show the loss of i-th iteration.

- **task**: Must be one of regression or classification. Based on task, this class encodes output values and also chooses activation function of last layer automatically.

This class implements this methods:

- **fit(x_train, y_train, validation: optional = list(x_valid, y_valid))**

- **predict(x)**

- **score(x,y)**

fit function returns the pair (history, validation_history) if the validation set was set. Otherwise it returns history of training phase. Each of these, contains a 2d numpy array with the shape of (n_iterations, 2). The columns contain loss and score history during training phase, respectively. predict function returns the predicted values for input x. score functions returns the score of model with input x and output y. This function uses accuracy_score and r2_score metrics implemented in scikit-learn module for classification and regression tasks.

7.2 CNN Model
For this case, we used the Keras API and implemented a simple class to be compatible with the previous MLP class.

```
CNN(architecture, epochs,
batch_size, optimizer, loss,
metrics, task, verbose)
```

where arguments are

- **architecture**: list of keras layers such as Cond2D, Dense, etc.

- **epochs**: number of epochs

- **batch_size**: size of the batch for gradient descent optimizer

- **optimizer**: keras optimizer such as adam or sgd

- **loss**: loss function, same as keras loss functions

- **metrics**: list of metrics, same as keras

- **task**: regression or classification

- **verbose**: print training information, same as keras

same as the MLP model, this class implements three methods fit, score and predict. everything is the same as the previous one except fit function which gets some keyword arguments and passes them to Keras fit function. So use Keras validation_data to validate the model on specific data.

7.3 RNN Model
Again, For this case, we used Keras API and implemented a simple class to be compatible with previous both MLP and CNN class.

```
RNN(architecture, epochs,
batch_size, optimizer, loss,
metrics, task, verbose)
```

where arguments are

- **architecture**: list of keras layers such as LSTM, Dense, etc.

- **epochs**: number of epochs

- **batch_size**: size of the batch for gradient descent optimizer

- **optimizer**: keras optimizer such as adam or sgd

- **loss**: loss function, same as keras loss functions

- **metrics**: list of metrics, same as keras

- **task**: regression or classification

- **verbose**: print training information, same as keras

same as the CNN model, this class implements three methods fit, score and predict.

7.4 Hyperparameter tuning
To find the optimal architecture of the network, grid search is a simple and efficient way. By passing a set of different possible choices for a problem, the following classes run and save all possible architectures and their corresponding loss history.

```
MLPGridSearch(task, hidden_layers,
activations batch_sizes,
epochs, mus, betas, etas,
alphas, filename)
```

for CNN

```
CNNGridSearch(architecture,epochs,
batch_sizes, optimizers,
file_name)
```

and RNN

```
RNNGridSearch(architecture,epochs,
batch_sizes, optimizers,
file_name)
```

These classes implement two methods run and best_model which return history of models and the best model which has the best accuracy and lower loss, respectively. Also, a parallelization technique for multi-core processors has been implemented on MLPGridSearch class to speed up the searching process. Because of GPU limitations, we cannot implement parallelization on CNNGridSearch and RNNGridSearch.

8. Results and Discussion

8.1 Iris
Dataset description:

- Input data: 4 features of 150 iris flowers

- Target value: 3 different types of iris

- More information in UCI repository

The grid search parameters were set to:

```
hidden_layers = [(5, 5, 5, 5),
(10, 10, 5),(15, 15),(20, 15, 10),
(32,)]
activations = [Tanh(), ReLu(),
LeakyReLu(.1)]
batch_sizes = [16, 64, 128]
epochs = [300]
mus = [0, .8]
betas = [.3]
etas = [.01, 0.001]
alphas = [.001, .01, 0]
```

See the results in the table 1. In this example, we saw that:

- LeakyReLu acts like Relu function in some situations.

- Maybe test accuracy becomes better that train accuracy.

- Momentum technique plays an important role in the optimization path. All of 6 best architectures are using momentum.

- Choosing learning is tricky and helps us to reach a global minimum in fewer epochs. By choosing big values for the learning rate, the training phase becomes unstable and the MLP class raises floating-point overflow.

- Smaller batch size helps us to reach a minimum faster.

8.2 Statlog

The database consists of the multi-spectral values of pixels in 3x3 neighborhoods in a satellite image and the classification associated with the central pixel in each neighborhood. The aim is to predict this classification, given the multi-spectral values. In the sample database, the class of a pixel is coded as a number. More information about this dataset is available in UCI repository.

The grid search parameters are set to

```
hidden_layers = [(32, 16, 8),
(10, 10, 10), (64,)]
activations = [Tanh(), Sigmoid(),
ReLu()]
batch_sizes = [256]
epochs = [30]
mus = [0.95]
betas = [.2, .3]
etas = [.01]
alphas = [0.01, 0]
```

See the results in table 2. By using these values, again we trained a network with this specification and reach train accuracy **83.74%** and test accuracy **81.45%** in 100 epochs.

```
MLP([64], activation=ReLu(),
batch_size=128,epochs=100,mu=0.95,
beta=.3, eta=.02, alpha=.001,
verbose=1, task='classification')
```


Figure 13. Loss and accuracy during epochs for train and validation sets of Statlog.

8.3 mnist

The MNIST database of handwritten digits, available from this page, has a training set of 60,000 examples, and a test set of 10,000 examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image. More information about mnist dataset is available on lecun.com.

The grid search parameters for this problem are set to

```
hidden_layers = [(64,64),(128,64),
(128, 32, 32),(128, 32, 32)]
activation=[ReLu(),LeakyReLu(.03)]
batch_size       s = [512]
epochs = [10]
mus = [0, .85]
betas = [.1,.2]
etas = [.001,.01]
alphas = [.001,0]
```

See the results in the table 3. Again we trained a network with this architecture and reached the **87.91%** accuracy on train set and **86.89%** on test set in just 15 epochs.

```
MLP([128, 64], activation=ReLu(),
batch_size=64, epochs=15, mu=0.85,
beta=.1, eta=.03, alpha=.001,
verbose=1, task='classification')
```

To reach this accuracy with just 15 epochs we reduced the batch size from 512 to 64. This shows the important role of batch size in training a network. Another interesting fact we found is the effect of the number of epochs on the overfitting

**Figure 12.** Loss value of best architectures for IRIS


Figure 14. Loss value of best architectures for Statlog

hidden_layer	activation	epoch	eta	beta	alpha	mu	batch_size	test_score	train_score	loss
(10, 10, 5)	ReLu	300	0.01	0.3	0.01	0.8	16	1.00	0.97	0.84
(10, 10, 5)	LeakyReLu(0.1)	300	0.01	0.3	0.01	0.8	16	1.00	0.96	1.28
(15, 15)	Tanh	300	0.01	0.3	0.01	0.8	16	0.97	0.98	0.67
(5, 5, 5, 5)	ReLu	300	0.01	0.3	0.01	0.8	16	0.97	0.97	0.85
(15, 15)	ReLu	300	0.01	0.3	0.01	0.8	16	0.97	0.97	0.87

Table 1. Best architectures and corresponding scores for IRIS dataset

hidden_layer	activation	epoch	eta	beta	alpha	mu	batch_size	test_score	train_score	loss
(64,)	ReLu	30	0.01	0.3	0.00	0.95	256	0.72	0.74	5.93
(64,)	Tanh	30	0.01	0.2	0.00	0.95	256	0.71	0.74	6.19
(64,)	Tanh	30	0.01	0.3	0.00	0.95	256	0.70	0.71	6.71
(64,)	ReLu	30	0.01	0.2	0.01	0.95	256	0.68	0.70	7.39
(64,)	Sigmoid	30	0.01	0.2	0.00	0.95	256	0.67	0.68	7.39
(64,)	ReLu	30	0.01	0.2	0.00	0.95	256	0.66	0.68	7.38
(64,)	ReLu	30	0.01	0.3	0.01	0.95	256	0.62	0.63	9.00

Table 2. Best architectures and corresponding scores for Statlog dataset

problem. In the above figure, you see that from iteration 7-8 the network validation accuracy is reducing.

Figure 16. Loss and accuracy during epochs for train and validation sets of MNIST.

8.4 Fashion mnist

Fashion-MNIST is a dataset of Zalando's article images consisting of a training set of 60,000 examples and a test set of 10,000 examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes. We intend Fashion-MNIST to serve as a direct drop-in replacement for the original MNIST dataset for benchmarking machine learning algorithms. It shares the same image size and structure of training and testing splits. More information in zalandoresearch's Github repository.

The grid search parameters are set to:

```
hidden_layers = [(64,64),(128,64),
(128, 32, 32),(128, 32, 32)]
activation=[ReLu(),LeakyReLu(.03)]
```

```
batch_sizes = [512]
epochs = [10]
mus = [0, .85]
betas = [.1,.2]
etas = [.001,.01]
alphas = [.001,0]
```

Again we trained a network with this architecture and reached the **76.78%** accuracy on train set and **75.36%** on test set in just 10 epochs.

```
MLP([128, 64], mu=0.85,
activation=LeakyReLu(.03),
batch_size=64, epochs=10,
beta=.2, eta=.01, alpha=.001,
verbose=1, task='classification')
```


Figure 18. Loss and accuracy during epochs for train and validation sets of Fashion MNIST.

hidden_layer	activation	epoch	eta	beta	alpha	mu	batch_size	test_score	train_score	loss
(128, 64)	ReLu	10	0.01	0.2	0.000	0.85	512	0.32	0.31	15.93
(128, 64)	LeakyReLu(0.03)	10	0.01	0.2	0.000	0.85	512	0.29	0.28	16.58
(128,)	ReLu	10	0.01	0.1	0.001	0.85	512	0.28	0.28	16.62
(128,)	ReLu	10	0.01	0.1	0.000	0.85	512	0.28	0.27	16.73
(64, 64)	ReLu	10	0.01	0.2	0.001	0.85	512	0.28	0.27	16.79
(128,)	LeakyReLu(0.03)	10	0.01	0.1	0.000	0.85	512	0.27	0.28	16.67

Table 3. Best architectures and corresponding scores for MNIST dataset

Figure 15. Loss value of best architectures for MNIST

hidden_layer	activation	epoch	eta	beta	alpha	mu	batch_size	test_score	train_score	loss
(128, 64)	LeakyReLu(0.03)	10	0.01	0.2	0.001	0.85	512	0.48	0.48	12.04
(128, 64)	ReLu	10	0.01	0.2	0.001	0.85	512	0.47	0.47	12.33
(128, 64)	ReLu	10	0.01	0.2	0.000	0.85	512	0.47	0.47	12.22
(128, 64)	LeakyReLu(0.03)	10	0.01	0.2	0.000	0.85	512	0.45	0.45	12.60
(64, 64)	LeakyReLu(0.03)	10	0.01	0.2	0.001	0.85	512	0.42	0.42	13.48
(128, 32, 32)	ReLu	10	0.01	0.2	0.000	0.85	512	0.39	0.40	13.84

Table 4. Best architectures and corresponding scores for Fashion MNIST dataset

Figure 17. Loss value of best architectures for Fashion mnist

The above figures shows that, unlike iris dataset, the validation accuracy is always lower that training accuracy.

To increase the accuracy, we changed the model from MLP to CNN and tested different architectures on this dataset.

Test 1) For the first test case we used a single convolution layer connected to a fully connected layer with 10 neurons and softmax activation function which classifies the input of the network. To show the effect of the convolution layer hyperparameters, we evaluated the network on five different feature maps $\{8, 16, 32, 64, 128\}$, five different kernel sizes $\{3 \times 3, 5 \times 5, 7 \times 7, 9 \times 9, 11 \times 11\}$ and two different padding strategies $\{same, valid\}$ for grid search parameters with Relu activation function, Keras default Adam optimizer, batch size 256 and 10 epochs. Table 6 shows the best result of these models. Although different kernel sizes and the number of filters are used to reach better accuracy, all of the results used only *same* padding.

Test 2) In the table 7 we extended this architecture with 2 convolutional layers. The results show that by increasing the network depth, 3×3 kernel size performs better than bigger kernels.

Test 3) In this experiment we increased the convolutional layers and epochs to 3 and 20, respectively. Also, we reduced batch size to 128. See table 8. In this table, even the train accuracy reached 100% but the test score is the same as previous examples. To overcome this issue we added a Dense layer with 32 neurons after the last convolutional layer. See table 9. In table 10 we changed filter sizes. These tables show

that increasing the number of filters instead of setting a fixed feature map, can perform better. All of the examples used smaller kernel sizes. So for the next examples, we use 3×3 kernel sizes along with increasing filter sizes trough depth of the network.

8.5 Cifar10

The CIFAR-10 dataset contains 60,000 32x32 color images in 10 different classes. The 10 different classes represent airplanes, cars, birds, cats, deer, dogs, frogs, horses, ships, and trucks. There are 6,000 images of each class. To show the effect of batch normalization effect on the dataset, we designed this architecture for cifar10 classification

```
[Conv2D(32, (3, 3),padding='same',
activation='relu'),
Conv2D(64, (3, 3),padding='same',
activation='relu'),
Conv2D(128, (3, 3),padding='same',
activation='relu'),
Flatten(),
Dense(256, activation='relu'),
Dense(128, activation='relu'),
Dense(10, activation='softmax')]
```

and ran the model with 10 epochs of adam optimizer and batch size 128. Without the normalization the accuracy of model on test set was 10% but when we added this layer, the accuracy became 64.4%. This model has $33,682,134$

parameters which should be trained using adam optimizer. Here by adding three pooling layers we construct this model

```
[Conv2D(32, (3, 3),padding='same',
activation='relu'),
MaxPooling2D((2,2)),
Conv2D(64, (3, 3),padding='same',
activation='relu'),
MaxPooling2D((2,2)),
Conv2D(128, (3, 3),padding='same',
activation='relu'),
MaxPooling2D((2,2)),
Flatten(),
Dense(256, activation='relu'),
Dense(128, activation='relu'),
Dense(10, activation='softmax')]
```

Although this model has just $651,990$ parameters the accuracy became 74.6% on the test set and 93% on train set which is a big achievement in both accuracy and speed. In this example we saw a little overfitting on the model since the difference between accuracy is about 18%. Here by add two dropout layers we reduced overfitting to about 8%. The new architecture is as follows

```
[Conv2D(32, (3, 3),padding='same',
activation='relu'),
Dropout(.5),
MaxPooling2D((2,2)),
Conv2D(64, (3, 3),padding='same',
activation='relu'),
MaxPooling2D((2,2)),
Conv2D(128, (3, 3),padding='same',
activation='relu'),
MaxPooling2D((2,2)),
Flatten(),
Dense(256, activation='relu'),
Dropout(.5),
Dense(128, activation='relu'),
Dense(10, activation='softmax')]
```

To show the effect of optimizer algorithm in the classification task when all of the other hyperparameters are fixed, we changed the optimizer from adam to RMSProp and SGD and saw that the accuracy decreased from 72.5% to 72.2 and 60.5%, respectively. For the last try, we changed the learning rate parameter of Adam to show the effect of the learning rate on the accuracy. See results in the table 5.

learning rate	epochs	accuracy
1e-1	10	10%
1e-2	10	42%
1e-3	10	74%
1e-4	10	64%
1e-1	50	10%
1e-2	50	47.4%
1e-3	50	74.8%
1e-4	50	72.8%

Table 5. The effect of epochs and learning for adam optimizer on cifar10 test set

8.6 Sentiment analysis

See dataset info in kaggle. We split the data set into 67% train and 33% test set. The first test case in this example consists of 108 tries, with 12 different architectures and 9 different permutations of sentence configuration. To do this we first preprocess the documents and convert sentences to a sequence of integers. To apply this sequence into the network we should align all sequences to a fixed size length. We tested 3 different lengths $100, 200, 500$. Since some words in the documents appear frequently we removed 15 top common words from the sequences. Then fed this fixed size sequences to an embedding layer with $10, 50, 100, 200$ output dimension along with one LSTM cell of $8, 16, 64$ units. Table 11 shows the accuracy of these models. All of the results obtained by just 5 epochs of Keras default Adam optimizer with batch size 256. As the second try, we designed a more complex network consist of an embedding layer followed by two LSTM cells. Table 12 shows the results of this experiment ran in 15 epochs. Also in this example, we used Bidirectional LSTM architecture to see the effect of this layer but doesn't affect the accuracy. For the last try, we designed a convolutional-LSTM network. This model performs better than previous models and reached 82% accuracy. Note that the first layer of this model is an embedding layer with output_dim=200, input_dim=715 and input_length=100. Here is the code of proposed architecture in the presented library.

```
N=max_dict_words + skip_common
[Embedding(N, 200, input_length=
max_sent_words),
LSTM(16, return_sequences=True),
Conv1D(16, 3, padding='valid',
activation='relu', strides=1),
GlobalMaxPooling1D(),
Dense(16,activation='relu'),
Dense(2, activation='softmax')]
```

9. Concluding remarks

In this exercise, we implemented a simple multi-layer perceptron neural network from scratch with matrix form formulations which are very important for fast prediction where

the number of neurons and layers is large. Our implementation is appropriate for both regression and classification tasks by just setting an argument. All of the preprocessing requirements such as scaling and encoding are done by *MLP* class automatically. For fine-tuning of hyperparameters of the network, we implemented a *MLPGridSearch* class which uses *parallelization* techniques and finds optimal architecture among all possible solutions. In the second part of the report, we implemented a *CNN* class based on an awesome *keras* module which is based on *tensor flow* a scientific computation package. This class has the same signature as the *MLP* class. Also to find optimal architecture, we implemented a *CNNGridSearch* class with the same interface of *MLPGridSearch* class. The third part implements *RNN* based networks with an interface similar to *CNN*. Again to optimize the network architecture, we implemented *RNNGridSearch*. Also to prevent the overfitting problem we added some regularization techniques to the models. We tested these networks on 6 different benchmark datasets which are used by many researchers for classification tasks. The results show the important role of model type and architecture, batch size, number of epochs, activation function and learning rate for training a network.

architecture	test	train	loss
Conv(filters=128, kernel=(11, 11))	0.88	0.93	0.19
Conv(filters=64, kernel=(5, 5))	0.87	0.95	0.15
Conv(filters=128, kernel=(5, 5))	0.87	0.95	0.14
Conv(filters=32, kernel=(7, 7))	0.87	0.94	0.16
Conv(filters=32, kernel=(9, 9))	0.87	0.93	0.18
Conv(filters=64, kernel=(9, 9))	0.87	0.93	0.20

Table 6. Train and test score for fashion mnist CNN with different configurations

architecture	test	train	loss
Conv(filters=32, kernel=(3, 3)),Conv(filters=32, kernel=(3, 3))	0.88	0.96	0.11
Conv(filters=32, kernel=(7, 7)),Conv(filters=32, kernel=(7, 7))	0.88	0.95	0.14
Conv(filters=16, kernel=(11, 11)),Conv(filters=16, kernel=(11, 11))	0.88	0.93	0.18
Conv(filters=128, kernel=(7, 7)),Conv(filters=128, kernel=(7, 7))	0.86	0.95	0.15
Conv(filters=64, kernel=(11, 11)),Conv(filters=64, kernel=(11, 11))	0.86	0.95	0.14
Conv(filters=16, kernel=(7, 7)),Conv(filters=16, kernel=(7, 7))	0.86	0.94	0.16

Table 7. Train and test score for fashion mnist CNN with different configurations

architecture	test	train	loss
Conv(filters=16, kernel=(3, 3)),Conv(filters=16, kernel=(3, 3)),Conv(filters=16, kernel=(3, 3))	0.84	0.99	0.04
Conv(filters=32, kernel=(7, 7)),Conv(filters=32, kernel=(7, 7)),Conv(filters=32, kernel=(7, 7))	0.82	0.98	0.04
Conv(filters=32, kernel=(7, 7)),Conv(filters=32, kernel=(7, 7)),Conv(filters=32, kernel=(7, 7))	0.82	0.98	0.06
Conv(filters=16, kernel=(3, 3)),Conv(filters=16, kernel=(3, 3)),Conv(filters=16, kernel=(3, 3))	0.80	0.99	0.03

Table 8. Train and test score for fashion mnist CNN with different configurations

architecture	test	train
Conv(filters=16, kernel=(3,3)),Conv(filters=16, kernel=(3, 3)),Conv(filters=16, kernel=(3,3)),Dense(units=32)	0.86	0.99
Conv(filters=16, kernel=(3,3)),Conv(filters=16, kernel=(3, 3)),Conv(filters=16, kernel=(3,3)),Dense(units=32)	0.85	1.00
Conv(filters=32, kernel=(7,7)),Conv(filters=32, kernel=(7, 7)),Conv(filters=32, kernel=(7,7)),Dense(units=32)	0.84	0.99
Conv(filters=32, kernel=(7,7)),Conv(filters=32, kernel=(7, 7)),Conv(filters=32, kernel=(7,7)),Dense(units=32)	0.83	0.99

Table 9. Train and test score for fashion mnist CNN with different configurations

architecture	test	train
Conv(filters=8, kernel=(3, 3)),Conv(filters=16, kernel=(3, 3)),Conv(filters=32, kernel=(3, 3)),Dense(units=32)	0.89	1.00
Conv(filters=8, kernel=(7, 7)),Conv(filters=16, kernel=(7, 7)),Conv(filters=32, kernel=(7, 7)),Dense(units=32)	0.86	0.99
Conv(filters=8, kernel=(3, 3)),Conv(filters=16, kernel=(3, 3)),Conv(filters=32, kernel=(3, 3)),Dense(units=32)	0.85	1.00
Conv(filters=8, kernel=(7, 7)),Conv(filters=16, kernel=(7, 7)),Conv(filters=32, kernel=(7, 7)),Dense(units=32)	0.85	0.99

Table 10. Train and test score for fashion mnist CNN with different configurations

architecture	test_score	train_score	loss
Embedding(output_dim=200, input_dim=515, input_length=100),LSTM(units=64)	0.814	0.840	0.35
Embedding(output_dim=200, input_dim=515, input_length=100),LSTM(units=8)	0.811	0.840	0.36
Embedding(output_dim=100, input_dim=515, input_length=100),LSTM(units=16)	0.809	0.840	0.35
Embedding(output_dim=200, input_dim=515, input_length=100),LSTM(units=16)	0.809	0.840	0.35
Embedding(output_dim=50, input_dim=515, input_length=100),LSTM(units=64)	0.806	0.830	0.39

Table 11. Train and test score for sentiment analysis with different configurations

architecture	test	train
Embedding(output_dim=200, input_dim=725, input_length=100),LSTM(units=32)	0.799	0.870
Embedding(output_dim=200, input_dim=725, input_length=100),LSTM(units=32)	0.796	0.920
Embedding(output_dim=500, input_dim=725, input_length=100),LSTM(units=32),LSTM(units=32)	0.795	0.950
Embedding(output_dim=1000, input_dim=725, input_length=100),LSTM(units=32)	0.795	0.930

Table 12. Train and test score for sentiment analysis with different configurations

architecture	test_score	train_score
LSTM(units=16),Conv1D(kernel=(3,), filters=16),GlobalMaxPooling1D(),Dense(units=16)	0.820	0.880
LSTM(units=16),Conv1D(kernel=(3,), filters=8),GlobalMaxPooling1D(),Dense(units=16)	0.818	0.880
LSTM(units=16),Conv1D(kernel=(5,), filters=8),GlobalMaxPooling1D(),Dense(units=16)	0.812	0.880
Conv1D(kernel=(3,), filters=16),GlobalMaxPooling1D(),Dense(units=16)	0.806	0.940

Table 13. Train and test score for sentiment analysis with different configurations

architecture	test_score	train_score	loss
Embedding(input_dim=625, input_length=100, output_dim=500),GRU(units=64)	0.814	0.860	0.33
Embedding(input_dim=625, input_length=100, output_dim=200),GRU(units=8)	0.814	0.850	0.35
Embedding(input_dim=625, input_length=100, output_dim=1000),GRU(units=64)	0.812	0.860	0.32
Embedding(input_dim=625, input_length=100, output_dim=200),GRU(units=16)	0.812	0.850	0.35
Embedding(input_dim=625, input_length=100, output_dim=500),GRU(units=16)	0.811	0.860	0.32

Table 14. Train and test score for sentiment analysis with different configurations

architecture	test	train
Embedding(output_dim=100, input_dim=325, input_length=300),GRU(units=8)	0.809	0.840
Embedding(output_dim=100, input_dim=325, input_length=300),GRU(units=16)	0.809	0.830
Embedding(output_dim=50, input_dim=325, input_length=300),GRU(units=64)	0.808	0.830
Embedding(output_dim=200, input_dim=325, input_length=300),GRU(units=16)	0.804	0.860
Embedding(output_dim=100, input_dim=325, input_length=300),GRU(units=64)	0.804	0.840
Embedding(output_dim=200, input_dim=325, input_length=300),GRU(units=64)	0.803	0.850
Embedding(output_dim=10, input_dim=325, input_length=300),GRU(units=64)	0.803	0.810
Embedding(output_dim=200, input_dim=325, input_length=300),GRU(units=8)	0.802	0.870

Table 15. Train and test score for sentiment analysis with different configurations

architecture	test	train
Embedding(input_dim=625, input_length=100, output_dim=10),GRU(units=64),Dense(units=8)	0.813	0.830
Embedding(input_dim=625, input_length=100, output_dim=500),GRU(units=8),Dense(units=8)	0.812	0.860
Embedding(input_dim=625, input_length=100, output_dim=500),GRU(units=64),Dense(units=8)	0.811	0.860
Embedding(input_dim=625, input_length=100, output_dim=10),GRU(units=8),Dense(units=8)	0.811	0.820
Embedding(input_dim=625, input_length=100, output_dim=500),GRU(units=16),Dense(units=8)	0.810	0.870

Table 16. Train and test score for sentiment analysis with different configurations